

CS2102 Helpsheet – AY2223 S1

Relational Algebra Equivalence

Conjunctive selection operations can be deconstructed into a sequence of individual selections; cascade of σ

$$\sigma_{\theta_1 \wedge \theta_2}(E) = \sigma_{\theta_1}(\sigma_{\theta_2}(E))$$

Only the final operation in a sequence of projection operations is needed, the rest can be omitted:

$$\Pi_{L_1}(\Pi_{L_2}(\dots(\Pi_{L_n}(E))\dots)) = \Pi_{L_1}(E)$$

Selections can be combined with Cartesian products and theta joins:

$$\sigma_{\theta}(E_1 \times E_2) = E_1 \bowtie_{\theta} E_2$$

$$\sigma_{\theta_1}(E_1) \bowtie_{\sigma_{\theta_2}} E_2 = E_1 \bowtie_{\sigma_{\theta_1 \wedge \theta_2}} E_2$$

Selection & Theta-join & Natural-join operations are commutative:

$$\sigma_{\theta_1}(\sigma_{\theta_2}(E)) = \sigma_{\theta_2}(\sigma_{\theta_1}(E))$$

$$E_1 \bowtie_{\theta} E_2 = E_2 \bowtie_{\theta} E_1$$

$$(E_1 \bowtie E_2) \bowtie E_3 = E_1 \bowtie (E_2 \bowtie E_3)$$

Theta joins are associative in the manner when θ_2 involves attributes from E_2 and E_3 only:

$$(E_1 \bowtie_{\theta_1} E_2) \bowtie_{\theta_2 \wedge \theta_3} E_3 = E_1 \bowtie_{\theta_1 \wedge \theta_3} (E_2 \bowtie_{\theta_2} E_3)$$

The selection operation distributes over the theta join operation under the following two conditions:

- It distributes when all the attributes in the selection condition σ_{θ_0} involves only the attributes of one of the expressions (E_1) being joined.

$$\sigma_{\theta_0}(E_1 \bowtie_{\theta} E_2) = (\sigma_{\theta_0}(E_1)) \bowtie_{\theta} E_2$$
- It distributes when the selection condition θ_1 involves only the attributes of E_1 and θ_2 involves only the attributes of E_2

$$\sigma_{\theta_1 \wedge \theta_2}(E_1 \bowtie_{\theta} E_2) = (\sigma_{\theta_1}(E_1)) \bowtie_{\theta} (\sigma_{\theta_2}(E_2))$$

The projection operation distributes over the theta join.

- Let L_1 and L_2 be attributes of E_1 and E_2 respectively. Suppose that the join condition θ involves only attributes in $L_1 \cup L_2$. Then

$$\Pi_{L_1 \cup L_2}(E_1 \bowtie_{\theta} E_2) = (\Pi_{L_1}(E_1)) \bowtie_{\theta} (\Pi_{L_2}(E_2))$$
- Consider a join $E_1 \bowtie_{\theta} E_2$. Let L_1 and L_2 be sets of attributes from E_1 and E_2 respectively. Let L_3 be attributes of E_1 that are involved in the join condition θ , but are not in $L_1 \cup L_2$, and let L_4 be attributes of E_2 that are involved in the join condition θ , but are not in $L_1 \cup L_2$. Then

$$\Pi_{L_1 \cup L_2}(E_1 \bowtie_{\theta} E_2) = \Pi_{L_1 \cup L_2}((\Pi_{L_1 \cup L_3}(E_1)) \bowtie_{\theta} (\Pi_{L_2 \cup L_4}(E_2)))$$

The set operations union and intersection (but not difference) are commutative.

$$E_1 \cup E_2 = E_2 \cup E_1$$

$$E_1 \cap E_2 = E_2 \cap E_1$$

Set union and intersection are associative.

$$(E_1 \cup E_2) \cup E_3 = E_1 \cup (E_2 \cup E_3)$$

$$(E_1 \cap E_2) \cap E_3 = E_1 \cap (E_2 \cap E_3)$$

The selection operation distributes over the union, intersection, and set-difference operations.

$$\sigma_P(E_1 - E_2) = \sigma_P(E_1) - \sigma_P(E_2) = \sigma_P(E_1) - \sigma_P(E_2)$$

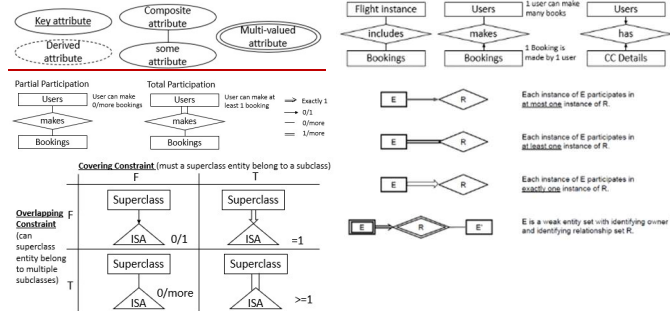
The projection operation distributes over the union operation.

$$\Pi_L(E_1 \cup E_2) = (\Pi_L(E_1)) \cup (\Pi_L(E_2))$$

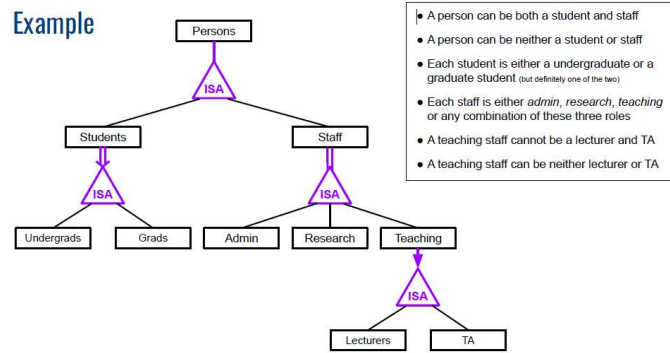
Relational Algebra Notes

Two relations R and S are **union-compatible** if R and S have the same attributes and corresponding attributes have the same or compatible domains, but R and S do not have to use the same attribute names.


ERD



Example



Aggregates



SQL

Possible Actions for ON DELETE and ON UPDATE:

- **NO ACTION:** Reject delete/update if violates constraint (DEFAULT)
- **RESTRICT:** Similar to NO ACTION except check cannot be deferred
- **CASCADE:** Propagates delete/update to referencing tuples
- **SET DEFAULT:** Updates FK of referencing tuples to some default value
- **SET NULL:** Updates FK of referencing tuples to NULL

Deferrable Constraints

Available for: UNIQUE, PRIMARY KEY, FOREIGN KEY

Example: `CONSTRAINT manager_fk FOREIGN KEY (manager) REFERENCES Employees (id) NOT DEFERRABLE`

→ **NOT DEFERRABLE, DEFERRABLE INITIALLY DEFERRED, DEFERRABLE INITIALLY IMMEDIATE**

In TX: `SET CONSTRAINTS ... DEFERRED`

Functions & Procedures

`CREATE OR REPLACE FUNCTION <name> (<param> <type>, <param> <type>, ...)`
`RETURN <type> AS $$`
`<CODE>`
`$$ LANGUAGE <language>`
 *where <language> includes sql or plpgsql

Control Structures

Variable


```
DECLARE [<var> <type>]*
<var> := <expr>
IF ... THEN ... [ELSIF ... THEN ...]* [ELSE ...] END IF
LOOP ... END LOOP
WHILE ... LOOP ... END LOOP
FOR ... IN ... LOOP ... END LOOP
BEGIN ... END
```

Selection
Repetition

Block

Cursor

Declare → Open → Fetch → Check (repeat) → Close



Cursor Movement

```
FETCH curs INTO r
FETCH PRIOR FROM curs INTO r
FETCH FIRST FROM curs INTO r
FETCH LAST FROM curs INTO r
FETCH ABSOLUTE n FROM curs INTO r
```

//from previous row
 //from top row
 //from last row
 //from nth row

Triggers

Trigger Functions

```
CREATE OR REPLACE FUNCTION <func>()
RETURNS TRIGGER AS $$
BEGIN
  <trigger_code>
END;
$$ LANGUAGE plpgsql;
```

Trigger Options

```
Events
• INSERT ON <table>
• DELETE ON <table>
• UPDATE [ OF <column> ] ON <table>

These set the TG_OP variable
```

Triggers

```
CREATE TRIGGER <name>
<trigger_timing>
ON <table>
FOR EACH <trigger_granularity>
[ WHEN <conditions> ]
EXECUTE FUNCTION <func>();
```

Timings

- AFTER/BEFORE (after or before the event)
- INSTEAD OF (replaces event, only for VIEWS)

Granularities

- FOR EACH ROW/FOR EACH STATEMENT

BEFORE trigger is executed before the EVENT

- May modify the EVENT by using return value
- May cancel the EVENT by using return value

AFTER trigger is executed after the EVENT

- The EVENT already occurred, o cancellation cannot be done using return value
- Can still be done using **RAISE EXCEPTION**

`TG_OP = {'INSERT', 'DELETE', 'UPDATE'}`

Triggers

Transition Variables

NEW: The modified row **after** the triggering event

OLD: The modified row **before** the triggering event

Events	NEW	OLD
INSERT	✓	✗
UPDATE	✓	✓
DELETE	✗	✓

Effect of Return Value

Events + Timings	NULL Tuple	NON-NULL Tuple t
BEFORE INSERT BEFORE UPDATE BEFORE DELETE	No insertion No update No deletion	Tuple t will be inserted/updated Deletion proceeds as normal (<i>t is ignored</i>)
AFTER INSERT AFTER UPDATE AFTER DELETE	NO EFFECT (<i>the operation is already done</i>)	

Trigger Granularities

Row-Level Trigger

FOR EACH ROW → Executes the trigger function for every tuple encountered

Statement-Level Trigger

FOR EACH STATEMENT → Executes the trigger for every statement regardless of the number of tuples

Timing	Row-Level	Statement-Level
AFTER	Tables	Tables and Views
BEFORE	Tables	Tables and Views
INSTEAD OF	Views	

Trigger Condition: WHEN

Example

```
CREATE TRIGGER bar
BEFORE INSERT ON Scores
FOR EACH ROW
WHEN (NEW.StuName = 'Adi')
EXECUTE FUNCTION foo();
```

Limitations

- **NO** SELECT in WHEN()
- **NO** OLD in WHEN() for INSERT
- **NO** NEW in WHEN() for DELETE
- **NO** WHEN() for INSTEAD OF

Deferred Trigger

Triggers that are checked only at the end of the TX instead of each statement

```
CREATE CONSTRAINT TRIGGER <trigger_name>
<trigger_timing> <trigger_event> ON <trigger_table>
[ DEFERRABLE INITIALLY [ DEFERRED | IMMEDIATE ] ]
FOR EACH <trigger_granularity>
[ WHEN <trigger_condition> ]
EXECUTE FUNCTION <trigger_function_name>();
```

Activation Order

For the same event on the same table:

1. BEFORE statement-level triggers
2. BEFORE row-level triggers
3. EVENT for the given ROW
4. AFTER row-level triggers
5. AFTER statement-level triggers

If (2) returns NULL, then (3) and (4) are cancelled
Within category: alphabetical order

Functional Dependencies

If we keep every attribute in one table and do not enforce the FD, we can experience anomalies:

- Redundant storage Update anomalies
- Deletion anomalies Insertion anomalies

The purpose of normal forms is to recognize designs that enforce functional dependencies by means of main SQL constraints and thus protect data against anomalies.

- A functional dependency $X \rightarrow Y$ is **trivial** iff $Y \subset X$.
- A functional dependency $X \rightarrow Y$ is **completely non-trivial** iff $Y \neq \emptyset$ and $Y \cap X = \emptyset$.
- Let R be a relation. Let $S \subset R$ be a set of attributes of R. S is a **superkey** of R iff $S \rightarrow R$.
- Let R be a relation. Let $S \subset R$ be a set of attributes of R. S is a **candidate** of R iff $S \rightarrow R$ and for all $T \subset S, T \neq S$, T is not a superkey of R.
- A **prime attribute** is an attribute that appears in some candidate key of R with Σ .
- Two sets of FD Σ and Σ' are equivalent iff they have the same closure: $\Sigma \equiv \Sigma', \Sigma^+ = \Sigma'^+$

Let Σ be a set of FDs of a relational schema R. The closure of a set of attributes $S \subset R$, denoted S^+ , is the set of all attributes that are functionally dependent on S. $S^+ = \{A \in R \mid \exists (S \rightarrow \{A\}) \in \Sigma^+\}$

Armstrong Axioms

Reflexivity: $\forall X \subset R \forall Y \subset R ((Y \subset X) \rightarrow (X \rightarrow Y))$

Augmentation: $\forall X \subset R \forall Y \subset R \forall Z \subset R ((X \rightarrow Y) \rightarrow (X \cup Z \rightarrow Y \cup Z))$

Transitivity: $\forall X \subset R \forall Y \subset R \forall Z \subset R ((X \rightarrow Y) \wedge (Y \rightarrow Z) \rightarrow (X \rightarrow Z))$

Other Axioms:

Weak Reflexivity: $\forall X \subset R (X \rightarrow \emptyset)$

Weak Augmentation: $\forall X, Y, Z \subset R ((X \rightarrow Y) \rightarrow (X \cup Z \rightarrow Y))$

Minimal Cover:

A set Σ of FD is **minimal** iff

1. The RHS of every FD in Σ is minimal. Every FD is of the form $X \rightarrow \{A\}$.
2. The LHS of every FD is minimal. Every FD in Σ of the form $X \rightarrow \{A\}$ there is no FD $Y \rightarrow \{A\}$ in Σ^+ s.t. Y is a proper subset of X.
3. The set itself is minimal. None of the FD in Σ can be derived from other FD in Σ .

A **minimal cover** of a set of FD Σ is set of FD Σ' that is both minimal and equivalent to Σ .

A set of FD is **compact** iff there is no different dependencies with same LHS. $\{\{A\} \rightarrow \{B\}, \{A\} \rightarrow \{C\}\}$ is not compact. $\{\{A\} \rightarrow \{B, C\}\}$ is compact.

A **compact cover** of a set of FDs Σ is set of FDs Σ' that is both compact and equivalent to Σ . A compact minimal cover of a set of FDs Σ is set of FDs Σ' that is both compact, minimal and equivalent to Σ .

BCNF

A relation R with a set of FD Σ is in BCNF $\leftrightarrow \forall \text{ FD } X \rightarrow \{A\} \in \Sigma^+:$

- $X \rightarrow \{A\}$ is trivial or
- X is a superkey

It is sufficient to look at Σ .

Decomposition

A **decomposition** of a table R is a set of tables $\{R_1, \dots, R_n\}$ s.t. $R = R_1 \cup \dots \cup R_n$.

A **binary decomposition** of a table R is a pair of tables $\{R_1, R_2\}$ s.t.

$R = R_1 \cup R_2$. A binary decomposition is **lossless-join** iff the full outer natural join of its two fragments equals the initial table.

Lossless-Join

A binary decomposition of R into R_1 and R_2 is **lossless-join** if $R = R_1 \cup R_2$ and

$R_1 \cap R_2 \rightarrow R_1$ or $R_1 \cap R_2 \rightarrow R_2$.

If $R_1 \cap R_2$ is the PK of one of the two tables, then it can be a FK in the other referencing the PK.

A decomposition is **lossless-join** if there exists a sequence of binary lossless-join decomposition that generates that decomposition.

Dependency Preserving

A decomposition of R with Σ into $R_1 \dots R_n$ with the respective sets of projected FD $\Sigma_1 \dots \Sigma_n$ is dependency preserving iff $\Sigma^+ = (\Sigma_1 \cup \dots \cup \Sigma_n)^+$.

Decomposition Algorithm

Let $X \rightarrow Y$ be a FD in Σ that violates the BCNF definition (it is not trivial and X is not a superkey). We use it to decompose R into the following two relations R_1 and R_2 .

- $R_1 = X^+$
- $R_2 = (R - X^+) \cup X$

We must now check whether R_1 and R_2 with the respective sets of projected FD Σ_1 and Σ_2 are in BCNF and continue the decomposition if they are not. The decomposition algorithm is guaranteed to find a **lossless decomposition in BCNF**, but may not be dependency preserving.

3NF

A relation R with a set of FDs Σ is in 3NF \leftrightarrow every FD $X \rightarrow \{A\} \in \Sigma^+$

- $X \rightarrow \{A\}$ is trivial or
- X is a superkey or
- A is a prime attribute

It is sufficient to look at Σ .

Bernstein Algorithm

When a relation is not in 3NF, we can synthesis a schema in 3NF from a **minimal cover** of the set of FDs.

- For each FD $X \rightarrow Y$ in the minimal cover create relation $R_i = X \cup Y$ unless it already exists or is subsumed by another relation.
- If none of the created relations contains one of the keys, pick a candidate key and create a relation with that candidate key.

In order to avoid unnecessary decomposition, it is generally a good idea to use a compact minimal cover. The algorithm guarantees a **lossless, dependency preserving decomposition in 3NF**.